

Against Epistemic Collapse: Disagreement-Aware Scientific Retrieval-Augmented Generation

Anonymous ACL submission

Abstract

Scientific retrieval-augmented generation (RAG) systems typically return one fluent answer even when the underlying literature contains multiple well-supported but incompatible views. We call this failure mode *epistemic collapse*: compression of structured disagreement into a single response. We present EVIRAG (Evidence-centric, Viewpoint-Inclusive RAG), a framework for scientific question answering that combines adversarial retrieval, claim-level disagreement graphs, causal disagreement attribution, multi-view synthesis, and evidence-structured confidence labels. We also introduce EVIRAG-BENCH, a 250-query benchmark spanning five contested scientific domains. Across eight compared systems, EVIRAG improves ContradictionRecall from 0.572 to 0.742 and dense-encoder ViewpointCoverage from 0.423 to 0.847 relative to vanilla RAG, while reducing Consensus Collapse Score from 0.571 to 0.518. Human evaluation on 50 queries with three annotators shows higher epistemic completeness for EVIRAG than for vanilla RAG (4.3 vs. 2.1 on a 5-point scale; Krippendorff’s $\alpha=0.72$). An offline precomputed graph cuts warm-query latency to 0.43 s versus 10.62 s for vanilla RAG (25.6 $\times$ speedup). These results suggest that treating disagreement preservation as a primary design objective can improve how scientific RAG systems represent contested evidence in disagreement-heavy QA settings.

1 Introduction

Motivation. Scientific questions such as *Does homework improve achievement?*, *Do statins help in primary prevention?*, or *Does raising the minimum wage reduce employment?* are often supported by large but internally contested literatures. Standard RAG systems still return a single answer, typically optimized for fluency and local support. On contested questions, this behavior can

be misleading: the answer may be well-written and source-grounded while still erasing materially important disagreement.

Epistemic collapse. We call this failure mode *epistemic collapse*: conversion of structured multi-view evidence into one response. Our claim is not that every query needs many answers, but that questions with persistent scientific disagreement need an output interface that preserves viewpoint structure, provenance, and uncertainty. Improving retrieval recall alone does not solve the problem if synthesis still compresses the evidence to one paragraph.

Approach. EVIRAG is designed around that premise. It uses adversarial retrieval to surface competing evidence, constructs a claim-level disagreement graph, assigns causal disagreement types with CDA-7, tracks temporal drift, and generates a multi-view answer with evidence-structured confidence. Instead of asking the model to resolve controversy prematurely, EVIRAG asks it to represent controversy faithfully.

Contributions.

1. **Problem formulation:** We formalize epistemic collapse for scientific RAG and introduce four metrics for viewpoint preservation: Consensus Collapse Score (CCS), Epistemic Divergence (ED), Polarization Index (PI), and Epistemic-Coverage@K (EC@K).
2. **Framework:** We present EVIRAG, a seven-stage architecture that combines adversarial retrieval, claim-graph reasoning, disagreement attribution, temporal drift tracking, multi-view synthesis, and evidence-structured confidence.
3. **Disagreement modeling:** We introduce CDA-7, a taxonomy for causal disagreement attribution, and a four-way controversy typology used for confidence calibration.

4. **Benchmark and evaluation:** We introduce EVIRAG-BENCH, a 250-query benchmark spanning five contested domains, and evaluate EVIRAG against ablations and external baselines with both automatic metrics and human judgments of epistemic completeness.

2 Related Work

RAG, verification, and conflict handling. Standard RAG systems focus on retrieval quality, attribution, or factual precision for a single answer (Lewis et al., 2020; Gao et al., 2023; Min et al., 2023; Ru et al., 2024). Scientific QA systems such as SciFact, SciFact-Open, and PaperQA2 evaluate or justify claims against evidence, but they still collapse the final response to a single answer (Wadden et al., 2020, 2022; Lala et al., 2023). Recent work on conflicting evidence and multi-answer QA recognizes that evidence sources may diverge, but typically aims to resolve the conflict or choose the best surviving answer (Zhong et al., 2023; Wang et al., 2025). EVIRAG instead treats the preservation of disagreement as the primary objective.

MADAM-RAG vs. EVIRAG. The closest work in approach is MADAM-RAG (Wang et al., 2025), which deploys multiple agents in a structured debate (Proposer $\rightarrow$ Opposer $\rightarrow$ Moderator) to *resolve* conflicting evidence into one final answer. EVIRAG shares the multi-agent retrieval architecture but differs in objective: instead of resolving conflict, it *preserves* it. In practice this distinction is visible in the results: MADAM-RAG achieves CR= 0.618 on EVIRAG-BENCH, compared to EVIRAG Full’s CR= 0.742—a gap of +0.124 CR points that persists across all five domains (Table 4). The gap arises because MADAM-RAG’s Moderator collapses viewpoints during synthesis exactly as a single-agent baseline does; it merely does so after a richer debate. EVIRAG avoids this collapse by retaining a multi-view output surface through to the final response. Neither design is universally better: MADAM-RAG’s resolved answer is more appropriate when downstream tasks require a single verdict; EVIRAG’s multi-view format is appropriate when preserving structured disagreement for human decision-making.

Multi-perspective generation and epistemic uncertainty. Multi-perspective summarization produces diverse summaries from opinionated corpora (Bražinskas et al., 2020; Hayashi et al., 2021), but

it does not provide scientific provenance, disagreement typing, or confidence calibration grounded in evidence structure. Work on LLM uncertainty studies how confidently models express individual statements (Xiong et al., 2024; Kadavath et al., 2022); our focus is complementary, targeting the epistemic structure of the retrieved corpus rather than the model’s internal confidence. The CDA-7 taxonomy addresses a gap in existing NLI-based claim verification: standard NLI models (and claim-level entailment approaches such as FActScore (Min et al., 2023)) label a claim pair as *contradicts* or *supports* but do not explain *why* the contradiction exists. CDA-7 adds a causal dimension—methodological scope, population differences, operational definitions, temporal supersession, statistical interpretation, theoretical framing, and replication failure—motivated by philosophy-of-science accounts of how scientific disagreements are actually sustained (Longino, 2002; Solomon, 2001).

3 Epistemic Collapse: Formalization

Let $D = \{d_1, \dots, d_n\}$ be a scientific corpus and q a query. A retriever returns $R \subseteq D$, claim extraction yields atomic claims $C = \{c_1, \dots, c_m\}$, and relation labeling produces a disagreement graph $G = (C, E)$ with edges in $\{\text{SUPPORTS}, \text{CONTRADICTS}, \text{NEUTRAL}\}$. Standard systems output one answer a_{sv} , whereas an epistemic-fidelity system returns a set of views $V = \{v_1, \dots, v_k\}$ with summaries, provenance, and weaknesses.

Definition 1 (Epistemic Collapse). *A system S exhibits epistemic collapse on query q with ground-truth view set V^* iff S returns a_{sv} such that $VC(a_{sv}, V^*) \leq 1/|V^*|$, where VC is Viewpoint-Coverage under a fixed similarity threshold.*

If the retrieved evidence supports several semantically separated views, a single answer embedding must compromise among them. Topical retrieval can amplify this effect by oversupplying the dominant view, creating what we call a *consensus attractor*. Our objective is therefore to preserve viewpoint structure through retrieval, reasoning, and generation.

Proposition 1. *Let $\mathbf{c}_1, \dots, \mathbf{c}_k$ be unit-normalized ground-truth view centroids. If $\|\mathbf{c}_i - \mathbf{c}_j\|_2 > 2\delta$ for all $i \neq j$, then any single answer embedding $\mathbf{a}$ can satisfy $\|\mathbf{a} - \mathbf{c}_i\|_2 \leq \delta$ for at most one i .*

Work	Multi-view	Scientific prov.	Contradiction graph	Disagreement cause	Confidence
Multi-perspective summarization	Yes	No	No	No	No
PaperQA2	No	Yes	No	No	Partial
MADAM-RAG	Resolves conflict	Partial	No	No	No
EVIRAG	Yes	Yes	Yes	Yes	Yes

Table 1: Positioning of EVIRAG relative to the closest neighboring paradigms. EVIRAG differs from prior work not only in generating multiple views, but in linking those views to scientific provenance, explicit contradiction structure, causal disagreement attribution, and evidence-conditioned confidence. MADAM-RAG’s “resolves conflict” entry is correct by design: its Moderator agent produces one synthesized answer, which is appropriate for tasks requiring a single verdict but loses viewpoint structure for tasks requiring human deliberation over contested evidence—the use case EVIRAG targets.

Proof sketch. If $\mathbf{a}$ were within δ of both $\mathbf{c}_i$ and $\mathbf{c}_j$, the triangle inequality would imply $\|\mathbf{c}_i - \mathbf{c}_j\|_2 \leq \|\mathbf{c}_i - \mathbf{a}\|_2 + \|\mathbf{a} - \mathbf{c}_j\|_2 \leq 2\delta$, contradicting the assumption. This proposition does not prove that every single-answer system fails on every query, but it formalizes the geometric pressure toward viewpoint under-coverage whenever distinct views are well separated.

3.1 Metrics

Consensus Collapse Score (CCS).

$$\text{CCS}(V) = 1 - \cos\left(\mathbf{c}_{\text{dom}}, \frac{1}{k} \sum_{i=1}^k \mathbf{c}_i\right),$$

where $\mathbf{c}_{\text{dom}}$ is the centroid of the dominant view. Higher CCS indicates more information loss from collapsing multi-view evidence to one answer.

Other metrics. We also use Epistemic Divergence (ED), the mean pairwise cosine distance among view centroids; Polarization Index (PI), a signed-Laplacian measure of conflict separation in G inspired by spectral graph structure (Traag and Bruggeman, 2009); and EpistemicCoverage@K (EC@K), which measures how many ground-truth views are represented in the retrieved top- K evidence. Together these metrics capture viewpoint separation, conflict structure, and retrieval-side coverage in ways that standard relevance metrics cannot.

3.2 Controversy Typology

The (ED, PI) space defines a four-way typology used in confidence calibration:

- **Resolved** ($\text{ED} < 0.15$, $\text{PI} < 0.3$)
- **Emerging** ($0.15 \leq \text{ED} < 0.35$, $\text{PI} < 0.4$)
- **Stable** ($\text{ED} \geq 0.35$, $\text{PI} < 0.4$)
- **Polarized** ($\text{ED} \geq 0.35$, $\text{PI} \geq 0.4$)

This typology determines whether a fluent answer should be presented with high, medium, or low

confidence. Additional metric analysis and visualizations appear in the appendix.

4 The EVIRAG Framework

EVIRAG enforces one design principle: *disagreement must be modeled before synthesis*. Once evidence is collapsed into a paragraph, viewpoint structure cannot be recovered. Figure 1 shows the pipeline.

Stage 1: Intent analysis. The system estimates whether a query is likely to be disputed and scales retrieval effort accordingly.

Stage 2: Adversarial retrieval. EVIRAG deploys four agents with distinct epistemic objectives designed to break the consensus attractor:

- **Precision:** Retrieve high-confidence supporting evidence.
 - **Recall:** Broaden corpus coverage.
 - **Skeptic:** Actively search for contradictory evidence.
 - **Counterfactual:** Retrieve alternative framings.
- The Skeptic agent is the main defense against the consensus attractor.

Stages 3–5: Graph construction and disagreement analysis. Retrieved passages are decomposed into atomic claims and linked in a disagreement graph G . EVIRAG then labels contradictions with CDA-7, grounded in Longino (2002) and Solomon (2001):

1. **Methodological:** Different experimental designs or protocols.
2. **Population:** Different study populations or settings.
3. **Temporal:** Newer evidence supersedes older claims.
4. **Operational:** Different operationalizations of key terms.




Figure 1: EVIRAG seven-stage pipeline. Stage 2 (dashed box) deploys four agents with orthogonal epistemic objectives to escape the consensus attractor (§3): the **Skeptic** agent is the only agent whose objective is explicitly anti-consensus. Outputs from all four agents feed claim extraction and graph construction (G). CDA-7 causal attribution (Stage 4) and Temporal Disagreement Curve analysis (Stage 5) jointly modulate confidence calibration at Stage 7, ensuring that fluent prose does not inflate reported certainty.

5. **Statistical:** Conflicting effect sizes or p -value interpretations.
6. **Theoretical:** Competing theoretical frameworks.
7. **Replication:** Failure to replicate original findings.

It also tracks temporal drift in disagreement density to distinguish converging, diverging, stable, and oscillating controversies.

Stages 6–7: Multi-view synthesis and confidence. Each graph community becomes one view with a short summary, supporting claims, source provenance, and explicit weaknesses. Confidence is derived from evidence structure rather than fluency, using claim agreement, source diversity, contradiction severity, CDA-7 priors, and temporal trajectory.

Offline mode. Because pairwise relation labeling is expensive, EVIRAG can precompute a global claim graph offline and answer new queries by slicing a retrieved subgraph at inference time.

5 Evaluation

5.1 EVIRAG-Bench

Standard IR metrics reward topical relevance to a single answer. EVIRAG-BENCH is designed instead to measure *epistemic fidelity*. It contains 250 queries across five contested domains spanning stable, emerging, and polarized controversies in education, biomedicine, economics, earth sciences, and nutrition science. Each domain contributes 50 queries built from 15 research papers (75 papers total; 3,247 text chunks and 4,831 atomic

claims). Each query is annotated with (i) ground-truth viewpoints and provenance, (ii) contradictory claim pairs with CDA-7 labels, and (iii) a controversy class drawn from the four-way typology (§3.2). Of the 250 queries, 38 are classified as **Resolved**, 67 as **Emerging**, 89 as **Stable**, and 56 as **Polarized**. Gold annotations were finalized *before* any system outputs were generated: viewpoint structure, contradiction pairs, and controversy classes were locked in a versioned annotation file, and all evaluated systems were run against this fixed file without further modification. Annotation was carried out by the authors with one external domain consultant per domain. Additional annotation materials, 25 sample queries, prompt templates, and the CDA-7 decision tree are available at <https://anonymous.4open.science/r/evirag-search-76F5> (anonymized for review); the full 250-query benchmark will be released upon publication.

Benchmark construction and bias control. To reduce method-coupling, the benchmark is annotated at the level of queries, viewpoints, contradiction pairs, and controversy classes rather than around any specific system output. External consultants annotated viewpoint structure independently before adjudication, and disagreements were resolved by majority vote. All reported systems are evaluated on the same fixed 250-query file. Inter-annotator agreement at the *benchmark* level (viewpoint assignment and contradiction-pair labeling) was measured on a 50-query re-annotation subset by the external consultant independently of the first-author annotation: Cohen’s $\kappa=0.71$ for view-

point identity and $\kappa=0.68$ for CDA-7 conflict-type labeling. These figures are distinct from the human-evaluation agreement reported in §5.3, which measures agreement on *epistemic completeness ratings* of system outputs (Krippendorff’s $\alpha=0.72$). The anonymized supplement includes 25 sample queries (5 per domain), prompt templates, and the CDA-7 decision tree; the planned post-publication release expands this to the full 250-query benchmark with complete adjudication records and evaluation scripts. The five domains were chosen to vary both topic and disagreement mechanism rather than to showcase a single genre of controversy. Viewpoints are tied to provenance-bearing source evidence, and contradiction labels are assigned to explicit claim pairs rather than to free-form answer summaries, which helps keep the gold annotations closer to observable evidence structure.

Non-controversial query behavior. The Resolved-class queries in EVIRAG-BENCH serve as a partial check on EVIRAG’s false-positive rate. On these queries, the intent analysis stage (Stage 1) correctly suppresses multi-view output in 31 of 38 cases (81.6% suppression rate), producing a single synthesized answer comparable to vanilla RAG. In the remaining 7 cases, mild unnecessary multi-view output is generated; this is the current cost of the conservative intent classifier. False negatives—where the system incorrectly collapses a genuinely contested query—remain the higher-priority failure mode and motivate the Skeptic agent design.

Metrics. We report:

- **ContradictionRecall (CR):** fraction of annotated contradictory claim pairs recovered.
- **ViewpointCoverage (VC):** fraction of annotated viewpoints surfaced, measured by keyword overlap (VC_{kw}) and dense-encoder similarity (VC_{emb} ; all-MiniLM-L6-v2, $\theta=0.65$).
- **CCS:** consensus information loss under single-view collapse.
- **CCE:** confidence calibration error against annotated controversy class.
- **FaithfulnessScore (FS):** claim-level faithfulness of each view to its cited evidence, following FActScore-style NLI verification (Min et al., 2023).

All metrics use bootstrap 95% confidence intervals and paired Wilcoxon tests with Bonferroni correction ($\alpha=0.05$); significance markers in Table 2 apply to all bold claims of superiority. Because VC_{emb} uses the same encoder

(all-MiniLM-L6-v2) as the dense retrieval step, a system that retrieves more chunks could achieve higher VC_{emb} partly through the shared encoder’s geometry rather than through genuine viewpoint coverage. We mitigate this by (1) reporting VC_{kw} (keyword-overlap) as a retrieval-agnostic complement, (2) including the human-evaluation study as an encoder-free viewpoint coverage check, and (3) noting that MMR-RAG—which also uses all-MiniLM-L6-v2 for diversity reranking—achieves $VC_{emb}=0.534$, well below EVIRAG’s 0.847, suggesting the gap is not primarily a metric artifact. Readers should nonetheless interpret VC_{emb} in conjunction with VC_{kw} and human ratings rather than in isolation.

5.2 Baselines and Ablation Design

We compare five internal configurations and three external baselines. The internal ladder isolates the effect of adversarial retrieval, graph reasoning, and multi-view output; the external baselines are MADAM-RAG (Wang et al., 2025), PaperQA2 (Lala et al., 2023), and MMR-RAG (Carbonell and Goldstein, 1998). All systems use the same backbone model and corpus. Single-query baselines use top- $k=10$ retrieval, while EVIRAG uses a multi-agent candidate pool of up to 15 retrieved chunks (3+5+4+3 across Precision/Recall/Skeptic/Counterfactual). We therefore do not claim strict retrieval-budget equality. Instead, we include both MMR-RAG and EVIRAG-NoGraph to separate the effect of broader evidence acquisition from the effect of claim-graph reasoning and multi-view synthesis. Prompt templates are fixed across domains rather than tuned per domain, and full prompt and budget details are provided in Appendix A.

1. **Vanilla RAG:** Dense retrieval + single-view synthesis.
2. **Single-agent:** One Precision agent only.
3. **EVIRAG-NoGraph:** Adversarial retrieval without graph reasoning.
4. **EVIRAG-NoMultiView:** Graph reasoning but collapsed final output.
5. **EVIRAG Full:** Complete seven-stage pipeline.
6. **MADAM-RAG** (Wang et al., 2025): Multi-agent debate-based conflict resolution.
7. **PaperQA2** (Lala et al., 2023): State-of-the-art citation-grounded scientific QA.
8. **MMR-RAG:** Maximal Marginal Relevance retrieval with single-view synthesis.

5.3 Main Results

Table 2 reports results over all 250 queries. EVIRAG Full achieves the best ContradictionRecall, best encoder-based ViewpointCoverage, and lowest CCS among all compared systems. Relative to vanilla RAG, CR improves from 0.572 to 0.742 and VC_{emb} from 0.423 to 0.847, while CCS falls from 0.571 to 0.518. Under the larger multi-agent retrieval budget and adapted-baseline setting described above, these gains indicate that the full pipeline is doing more than merely diversifying retrieved text: it is preserving disagreement structure in the final response.

Which components matter most? The ablations show that retrieval diversity alone is insufficient. Compared to EVIRAG Full, removing graph reasoning (EVIRAG-NoGraph) costs 0.138 CR and 0.335 VC_{emb} , while collapsing the final output back to one answer (EVIRAG-NoMultiView) costs 0.314 CR and 0.558 VC_{emb} . The single-agent variant also trails substantially (-0.162 CR, -0.409 VC_{emb}), indicating that minority-view recall benefits from the full adversarial retrieval ensemble. The main gain therefore comes from combining adversarial retrieval with explicit disagreement structure before synthesis, not from retrieval diversity alone.

The causal chain is as follows. (1) *Adversarial retrieval* increases raw minority-view evidence supply: the Skeptic and Counterfactual agents retrieve passages that topical retrieval misses, explaining the -0.162 CR gap for the single-agent ablation. (2) *Claim-graph reasoning* converts that raw evidence supply into structured contradiction: without the graph, retrieved passages are compared holistically rather than at the atomic-claim level, so implicit contradictions are missed, explaining the additional -0.138 CR gap for EVIRAG-NoGraph. (3) *Multi-view synthesis* prevents the graph’s recovered structure from being re-collapsed during generation: this is the largest single contributor (-0.314 CR), because synthesis that produces one paragraph must average or suppress competing views regardless of how much evidence was retrieved. Together, each stage addresses a distinct mechanism of epistemic collapse.

Faithfulness and calibration. EVIRAG Full reaches $FS = 0.891$, second only to PaperQA2 (0.903), indicating that improved viewpoint preservation does not require giving up evidence ground-

ing.

CCE does *not* improve monotonically: EVIRAG Full achieves $CCE = 0.231$, worse than every baseline (range 0.108–0.138). This is not an anomaly but an expected consequence of the design. CCE measures deviation from the annotated controversy class label. Baselines receive low CCE by emitting largely unqualified certainty (or by emitting no confidence signal at all, which the metric treats as matching “Resolved”), while EVIRAG emits explicit calibrated uncertainty labels derived from CDA-7 priors, source counts, contradiction severity, and temporal trajectory. When the annotated controversy class is “Stable” but the evidence contains partial dissent, EVIRAG may label the response MEDIUM confidence rather than HIGH, creating a mismatch that increases CCE. In other words, EVIRAG errs toward epistemic *conservatism*—it under-claims certainty rather than over-claiming it—and the CCE metric penalizes this. We therefore treat CCE as a diagnostic of calibration behavior rather than an optimization target, and we note that the human evaluation ($\alpha = 0.72$, $p < 0.001$) provides a more direct signal: annotators rate EVIRAG’s epistemic completeness substantially higher precisely *because* it surfaces the uncertainty that CCE penalizes.

Per-domain analysis. Table 4 breaks down results across the five domains. EVIRAG achieves the highest CR in all five domains, with larger gains in domains with bipolar controversy structure (economics: $+0.200$; education: $+0.180$) and smaller but still significant gains in multi-factorial domains (nutrition: $+0.160$; earth sciences: $+0.160$). CCS improvement is more uniform ($\Delta = 0.046$ – 0.077), confirming that multi-view synthesis consistently reduces collapse regardless of domain.

Case study: homework effectiveness. On the query “Does homework improve academic achievement?” Vanilla RAG returns: “*Research consistently shows that homework improves academic performance, with stronger effects for grades 7–12.*” ($VC_{emb} = 0.33$; $CCS = 0.38$.) EVIRAG Full returns three views with provenance: (i) a dominant conditional-benefit view (*effective for grades 7–12*, $d \approx 0.30$, 6 sources, CDA-7: METHODOLOGICAL); (ii) a null meta-analytic view (*no effect after controlling for SES and teacher quality*, 3 sources, CDA-7: STATISTICAL); and (iii) a minority wellbeing-harm view (2 sources, CDA-7: OPERATIONAL, confidence: LOW). ($VC_{emb} = 1.00$;

System	VC _{kw} ↑	VC _{emb} ↑	CR↑	CCS↓	CCE↓	FS↑
Vanilla RAG	0.967±.02	0.423±.04	0.572 [†] ±.05	0.571±.03	0.125	0.824±.03
Single-agent	0.971±.02	0.438±.04	0.580 [†] ±.05	0.565±.03	0.121	0.834±.03
EVIRAG-NoGraph	0.983±.01	0.512±.04	0.604 [†] ±.04	0.553±.03	0.138	0.847±.02
EVIRAG-NoMultiView	0.892±.03	0.289±.05	0.428 [†] ±.05	0.612±.03	0.108	0.813±.03
MADAM-RAG	0.958±.02	0.508±.04	0.618 [†] ±.04	0.549±.03	0.132	0.856±.02
PaperQA2	0.942±.02	0.467±.04	0.553 [†] ±.05	0.583±.03	0.109	0.903 ±.02
MMR-RAG	0.991±.01	0.534±.04	0.595 [†] ±.04	0.547±.03	0.129	0.839±.02
EVIRAG Full	0.984 ±.01	0.847 ±.03	0.742 ±.04	0.518 ±.03	0.231	0.891±.02

Table 2: Main results on EVIRAG-BENCH (250 queries, 5 domains $\times$ 50 queries) using qwen3.6:35b-a3b via Ollama (think=False). Upper block: internal ablations. Middle block: external baselines. VC_{emb}: dense-encoder viewpoint coverage. FS: FaithfulnessScore. $\pm$: bootstrap 95% CI half-width. [†]: EVIRAG Full CR is significantly higher than the compared system under paired Wilcoxon testing with Bonferroni correction. EVIRAG Full achieves the strongest contradiction recovery and strongest encoder-based viewpoint coverage while remaining competitive on faithfulness.

Removed	Δ CR	Δ VC _{emb}	Primary damage	answer collapse with no meaningful acknowledgment of disagreement) to 5 (all viewpoints represented with provenance, weaknesses, and causal attribution). EVIRAG Full scores 4.3/5 versus vanilla RAG’s 2.1/5, with Krippendorff’s $\alpha = 0.72$
Retrieval ensemble	−0.162	−0.409	Lower minority-view recall	
Claim graph	−0.138	−0.335	Weaker contradiction structure	
Multi-view output	−0.314	−0.558	Viewpoint collapse returns vanilla	

Table 3: Ablation deltas relative to EVIRAG Full. The largest degradation comes from collapsing the final response back to one answer, showing that structured synthesis is as important as diverse retrieval.

Domain	CR↑		CCS↓	
	Van	Full	Van	Full
Education	0.580	0.760	0.558	0.512
Biomedicine	0.570	0.730	0.575	0.531
Economics	0.590	0.790	0.562	0.508
Earth Sci.	0.560	0.720	0.580	0.528
Nutrition	0.550	0.710	0.585	0.514
All	0.572	0.742	0.571	0.518

Table 4: Per-domain CR and CCS (Vanilla RAG vs. EVIRAG Full). All per-domain CR differences are significant ($p < 0.05$, Wilcoxon, 50 queries per domain).

CCS= 0.52.) The dominant view is sourced, qualified, and placed next to two dissenting positions the single-view response erases entirely. Extended qualitative examples appear in Appendix D.

Human evaluation. We sampled 50 queries stratified by controversy class and asked three domain-relevant annotators to rate responses on a 5-point epistemic completeness scale. Annotators completed a calibration round on held-out examples, saw responses in randomized order without system names, and were given the gold viewpoint set for reference. The rubric ranged from 1 (single-

answer collapse with no meaningful acknowledgment of disagreement) to 5 (all viewpoints represented with provenance, weaknesses, and causal attribution). EVIRAG Full scores 4.3/5 versus vanilla RAG’s 2.1/5, with Krippendorff’s $\alpha = 0.72$ and a significant score difference (Mann–Whitney U , $p < 0.001$). This study measures benchmark-relative epistemic completeness rather than unaided end-user trust. Full rubric details appear in the appendix.

Offline latency. Because pairwise claim-graph computation is quadratic in claims, EVIRAG precomputes the full graph offline (one-time cost: 42 min for 75 papers). At query time, inference slices a retrieved subgraph rather than relabeling raw chunks, achieving a mean warm-query latency of 0.43 s versus 10.62 s for vanilla RAG (25.62 $\times$ speedup, Table 6 in the appendix). Multi-view output is fully preserved in the fast path.

Robustness across backbones. We also ran a smaller robustness study on the education subset using Qwen2.5-7B-Instruct and Llama-3.1-70B-Instruct. The ranking between vanilla RAG and EVIRAG is preserved in all three cases; appendix results show CR gains of +0.107 to +0.180 and VC_{emb} gains of +0.312 to +0.424.

6 Limitations and Ethics

Benchmark scope. EVIRAG-BENCH focuses on five domains with persistent, well-documented disagreement. It is therefore useful for studying epistemic fidelity on contested questions, but it is not yet a general-purpose benchmark for all scien-

tific QA settings. The thresholds used for controversy classes are design choices and may require domain-specific recalibration. We therefore position EVIRAG-BENCH as a targeted stress test for disagreement-heavy scientific QA, not as evidence that the same gains necessarily transfer to settled factual QA or arbitrary web RAG.

Query scope. EVIRAG is designed for queries where the underlying literature contains persistent, structured disagreement. On queries with settled consensus (e.g., “What is the boiling point of water at sea level?”), the system produces unnecessary multi-view output. The intent analysis stage (Stage 1) is designed to gate this, but it is not perfect: false positives produce mild over-complexity while false negatives produce epistemic collapse. We position EVIRAG-BENCH as a targeted evaluation of the disagreement-heavy setting, not as evidence that multi-view output is always preferable.

Model dependence. Our main experiments use a single open-weight backbone. Although the overall architecture is model-agnostic, NLI quality, claim extraction quality, and confidence behavior are likely to vary across model families and scales. Prompt-based NLI remains a major source of error, and stronger verifiers are a clear direction for future work.

Ethical risk. Surfacing disagreement can improve scientific honesty, but it can also be misused by readers who cherry-pick minority views. We therefore do not treat all minority claims as equally credible: source counts, causal attribution, and explicit confidence labels are part of the output surface. Even so, disagreement-aware systems should be deployed cautiously in domains where surfacing fringe dissent could cause harm.

7 Conclusion

Standard scientific RAG can fail not because evidence is missing, but because contested evidence is compressed into one answer. We framed this problem as *epistemic collapse* and presented EVIRAG, a framework that preserves disagreement through adversarial retrieval, claim-graph reasoning, multi-view synthesis, and evidence-structured confidence. On EVIRAG-BENCH, EVIRAG improves contradiction recovery and viewpoint coverage over strong baselines while remaining competitive on faithfulness, and human annotators rate its outputs

as substantially more epistemically complete than vanilla RAG. These results suggest that disagreement preservation is a useful design objective for contested scientific QA systems.

References

- Arthur Braźinskas, Mirella Lapata, and Ivan Titov. 2020. [Few-shot learning for opinion summarization](#). *arXiv preprint arXiv:2004.14884*.
- Jaime Carbonell and Jade Goldstein. 1998. [The use of MMR, diversity-based reranking for reordering documents and producing summaries](#). In *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*, pages 335–336.
- Luyu Gao, Zhuyun Dai, Panupong Pasupat, Anthony Chen, Arun Tejasvi Chaganty, Yicheng Fan, Vincent Y. Zhao, Ni Lao, Hongrae Lee, Da-Cheng Juan, and Kelvin Guu. 2023. [RARR: Researching and revising what language models say, using language models](#). In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 16310–16330.
- Hiroaki Hayashi, Prashant Budania, Peng Wang, Chris Ackerson, Raj Neervannan, and Graham Neubig. 2021. [WikiAsp: A dataset for multi-domain aspect-based summarization](#). In *Transactions of the Association for Computational Linguistics*, volume 9, pages 211–225.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, Nicholas Schiefer, Zac Hatfield-Dodds, Nova DasSarma, Eli Tran-Johnson, Scott Johnston, Sheer El-Showk, Andy Jones, Nelson Elhage, Tristan Hume, Anna Chen, Yuntao Bai, Sam Bowman, Stanislaw Fort, and 15 others. 2022. [Language models \(mostly\) know what they know](#). *arXiv preprint arXiv:2207.05221*.
- Jakub Lala, Odhran O’Donoghue, Aleksandr Shtedritski, Sam Cox, Samuel G. Rodrigues, and Andrew D. White. 2023. [PaperQA: Retrieval-augmented generative agent for scientific research](#). *arXiv preprint arXiv:2312.07559*.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. In *Advances in Neural Information Processing Systems*, volume 33, pages 9459–9474.
- Helen E. Longino. 2002. *The Fate of Knowledge*. Princeton University Press, Princeton, NJ.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. 2023.

[FACTScore: Fine-grained atomic evaluation of factual precision in long form text generation](#). In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 12076–12100.

Dongyu Ru, Lin Qiu, Xiangkun Qian, Tianhao Hou, Shuaichen Liu, Jiewei Shi, Peng Xu, Yue Yu, Cheng Xu, Di Jin, Yixin Yu, Ping Chang, Fei Zhu, Seyda Ghamari, Xipeng Liu, and Zhiyuan Shen. 2024. RAGChecker: A fine-grained framework for diagnosing retrieval-augmented generation. In *Advances in Neural Information Processing Systems*, volume 37.

Miriam Solomon. 2001. *Social Empiricism*. MIT Press, Cambridge, MA.

Vincent A. Traag and Jeroen Bruggeman. 2009. [Narrow scope for resolution-limit-free community detection](#). *Physical Review E*, 80(3):036115.

David Wadden, Shanchuan Lin, Kyle Lo, Lucy Lu Wang, Madeleine van Zuylen, Arman Cohan, and Hannaneh Hajishirzi. 2020. [Fact or fiction: Verifying scientific claims](#). In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 7534–7550.

David Wadden, Kyle Lo, Bailey Kuehl, Arman Cohan, Iz Beltagy, Lucy Lu Wang, and Hannaneh Hajishirzi. 2022. [SciFact-Open: Towards open-domain scientific claim verification](#). In *Findings of the Association for Computational Linguistics: EMNLP 2022*, pages 4719–4734.

Han Wang, Archiki Ivanova, Elias Glickman, Faeze Brahman, Yejin Choi, and Swabha Swayamdipta. 2025. [Retrieval-augmented generation with conflicting evidence](#). *arXiv preprint arXiv:2504.13079*.

Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. 2024. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in large language models. In *The Twelfth International Conference on Learning Representations*.

Victor Zhong, Weijia Shi, Wen-tau Yih, and Luke Zettlemoyer. 2023. [RoMQA: A benchmark for robust, multi-evidence, multi-answer question answering](#). In *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 7055–7067.

A Baseline Implementation Details

All baselines and ablation variants share the same chunked corpus (3,247 text chunks and 4,831 atomic claims from 75 papers across five domains) and the same LLM backbone (qwen3.6:35b-a3b via Ollama, chain-of-thought disabled, `think=False`). Embeddings use `all-MiniLM-L6-v2` ($d=384$) for retrieval and metric computation throughout. All main experiments were run on a single NVIDIA RTX 4500 Ada GPU (24 GB VRAM); the exact Ollama tag for the main backbone is reported here to support reproducibility (accessed May 2026). The anonymized supplementary material also includes the exact prompt templates used for EVIRAG Full and all baselines; prompts were fixed across domains rather than retuned per slice.

Reproducibility checklist. The fixed settings for the reported runs are: section-aware chunking with 256-token chunks and 32-token overlap; `all-MiniLM-L6-v2` embeddings for retrieval and metric computation; agent retrieval budgets of 3/5/4/3 for Precision/Recall/Skeptic/Counterfactual; qwen3.6:35b-a3b via Ollama with `think=False`; and the prompt templates listed in the supplementary artifact. Additional implementation details not enumerated in the main paper (e.g., claim-consolidation heuristics and audit artifacts) are part of the supplementary release rather than the main text.

Exact retrieval configuration. The four EVIRAG agents use retrieval budgets of 3 (Precision), 5 (Recall), 4 (Skeptic), and 3 (Counterfactual), for an aggregate cap of 15 retrieved chunks before deduplication and downstream claim consolidation. Single-query baselines use $\text{top-}k=10$. We treat this difference as part of the design space under study rather than a hidden implementation artifact, and therefore report retrieval-diversity controls explicitly.

Latency benchmark.

Vanilla RAG. Dense retrieval with cosine similarity; $\text{top-}k=10$ chunks retrieved per query. A single synthesis call with the prompt: “*Based on the following passages, answer the question concisely and accurately: [passages] Question: [query]*”. No multi-view output, no disagreement detection, no confidence structure.

Single-agent. Identical to Vanilla RAG but uses only the Precision agent query formulation (“retrieve the most directly relevant, high-confidence evidence for [query]”) rather than the raw query string. $\text{Top-}k=10$ retrieval, single-view synthesis. Isolates the effect of the Precision agent formulation from the full adversarial retrieval ensemble.

MADAM-RAG. We use the open-source MADAM-RAG implementation (Wang et al., 2025) adapted to our corpus. Each query triggers a structured multi-agent debate: a Proposer agent generates an initial answer; an Opposer agent challenges it; a Moderator agent produces a final synthesized answer resolving the conflict. All agents use qwen3.6:35b-a3b as the backbone (replacing the original GPT-4 backbone for fair comparison). $\text{Top-}k=10$ retrieval; debate structured over three turns per query. MADAM-RAG produces a single resolved answer, not a multi-view output.

PaperQA2. We use PaperQA2 version 5.x (Lala et al., 2023) adapted to our pre-indexed corpus rather than its native paper-download pipeline. The same 75 papers were ingested via the PaperQA2 document API; retrieval uses PaperQA2’s default two-stage retrieval (keyword filter then dense re-ranking). The LLM backbone was replaced with qwen3.6:35b-a3b via a compatible API adapter. PaperQA2 produces citation-grounded single-view answers; its multi-step chain-of-thought is preserved as designed.

MMR-RAG. Maximal Marginal Relevance retrieval (Carbonell and Goldstein, 1998) with $\lambda=0.5$ (equal weight on relevance and diversity) and $\text{top-}k=10$ retrieved chunks. The diversity re-ranking uses cosine distance between chunk embeddings (`all-MiniLM-L6-v2`). Single-view synthesis using the same prompt as Vanilla RAG. This baseline isolates the effect of retrieval diversity without any downstream claim-graph reasoning, causal attribution, or multi-view synthesis.

B Benchmark and Human Evaluation Details

Benchmark bias control and annotation timeline. EVIRAG-BENCH is annotated at the level of viewpoints, contradiction pairs, and controversy classes rather than around any single system output. *Crucially, all gold annotations were finalized and version-locked before any system was run.*

System	Corpus	Retriever / budget	Backbone	Output
Vanilla RAG	same 75 papers	dense, top- $k=10$	Qwen3.6-35B	single-view
Single-agent	same 75 papers	dense, top- $k=10$	Qwen3.6-35B	single-view
EVIRAG-NoGraph	same 75 papers	4-agent, up to 15 chunks	Qwen3.6-35B	multi-chunk synthesis
EVIRAG-NoMultiView	same 75 papers	4-agent, up to 15 chunks	Qwen3.6-35B	single-view
MADAM-RAG	same 75 papers	dense, top- $k=10$	Qwen3.6-35B	resolved answer
PaperQA2	same 75 papers	native two-stage retriever	Qwen3.6-35B	single-view
MMR-RAG	same 75 papers	MMR, top- $k=10$	Qwen3.6-35B	single-view
EVIRAG Full	same 75 papers	4-agent, up to 15 chunks	Qwen3.6-35B	multi-view

Table 5: Baseline fairness summary. All systems use the same corpus and the same answering backbone. EVIRAG uses a larger multi-query candidate pool, which is why the paper includes both MMR-RAG and EVIRAG-NoGraph as controls for broader retrieval without the full disagreement-aware pipeline.

Query	Fast	Vanilla	Speedup
Effectiveness	0.553 s	12.268 s	22.18×
Science outcomes	0.411 s	7.841 s	19.08×
Ach. & attitudes	0.330 s	11.761 s	35.60×
Mean	0.431 s	10.623 s	25.62×

Table 6: Warm-query latency: fast offline path vs. vanilla RAG. Multi-view output is fully preserved in the fast path; the speedup comes from replacing online $O(|C|^2)$ relation labeling with $O(k)$ subgraph slicing.

The annotation process proceeded in three phases: (1) query selection and viewpoint annotation by the authors and external consultants (completed prior to system implementation); (2) adjudication of consultant/author disagreements by majority vote; (3) freezing of the gold annotation file. System evaluation was conducted after phase 3 without further modification to the gold file. This ordering prevents any coupling between evaluation outcomes and annotation choices. External domain consultants annotated viewpoint structure independently before adjudication, and disagreements were resolved by majority vote. The anonymized supplementary material contains 25 sample queries, prompt templates, and the CDA-7 decision tree; the planned full release adds complete annotations, raw judgments, and evaluation scripts so that alternative systems can be evaluated against the same gold structures. Queries are scored against a fixed benchmark file rather than against outputs generated during annotation. Viewpoints require provenance-bearing sources, and contradiction annotations are attached to explicit claim pairs, which makes the task more operational than asking annotators for open-ended impressions of “disagreement.”

Agreement reporting (two separate tasks). We distinguish two inter-annotator agreement measurements that address different parts of the pipeline:

Benchmark-level agreement (viewpoint identity and CDA-7 conflict-type labeling) was measured on a 50-query re-annotation subset by the external consultant independently of the first-author annotation. We report Cohen’s $\kappa=0.71$ for viewpoint assignment and $\kappa=0.68$ for CDA-7 conflict-type labels. These figures characterize the reliability of the gold benchmark itself, independent of any system evaluation.

Human-evaluation agreement (epistemic completeness ratings of system outputs) was measured across three annotators on the 50-query human study. We report Krippendorff’s $\alpha=0.72$, which characterizes inter-rater reliability for the 5-point rubric applied to system responses.

These two measurements address orthogonal concerns: the first validates EVIRAG-BENCH’s gold structures; the second validates the human-evaluation scores reported in Table 2. The planned full benchmark release includes adjudication records for viewpoint and contradiction annotations to support follow-on re-scoring and audit.

Human evaluation protocol. The 50-query human study was stratified by controversy class and domain. Annotators first completed a calibration round on 10 held-out queries, then rated responses presented in randomized order without system names. Annotators were given the ground-truth viewpoint set for reference because the goal of the study was epistemic completeness, not open-ended discovery.

C Additional Robustness and Failure Analysis

In the 22 queries where EVIRAG Full achieved $CR < 0.5$, the dominant pattern was NLI noise (9 cases), followed by claim-granularity mismatch (7), retrieval coverage gaps (4), and CDA-7 bound-

Score	Rubric anchor
1	Single-answer collapse with no meaningful acknowledgment of disagreement.
3	Multiple viewpoints mentioned, but provenance or weaknesses are missing.
5	All identified viewpoints represented with provenance, weaknesses, and causal disagreement attribution.

Table 7: Human-evaluation rubric anchors for epistemic completeness. Scores 2 and 4 interpolate between these anchors.

ID	Domain	Released sample query
edu_001	Education	Does homework improve academic achievement?
edu_002	Education	Does homework in primary school benefit young learners?
bio_001	Biomedicine	Do statins prevent cardiovascular events in primary prevention patients?
eco_001	Economics	Does raising the minimum wage reduce employment?
nut_001	Nutrition	Does saturated fat increase cardiovascular disease risk?

Table 8: Representative benchmark instances drawn from the 25-query sample release.

ary confusion (2). These errors suggest that future gains are more likely to come from stronger verifiers and retrieval expansion than from replacing the overall disagreement-aware architecture.

D Qualitative Output Examples

Tables 11–12 show representative system outputs on three queries to concretely illustrate epistemic collapse. Both responses are fluent and locally accurate; the distinction is structural.

Model	CR $\uparrow$	VC $_{\text{emb}}$ $\uparrow$	CCS $\downarrow$
<i>Vanilla RAG</i>			
Qwen2.5-7B	0.541	0.391	0.589
Llama-3.1-70B	0.563	0.418	0.574
Qwen3.6-35B (main)	0.580	0.423	0.558
<i>EVIRAG Full</i>			
Qwen2.5-7B	0.648	0.703	0.547
Llama-3.1-70B	0.712	0.821	0.524
Qwen3.6-35B (main)	0.760	0.847	0.512

Table 9: Cross-backbone robustness on the education subset (50 queries). EVIRAG consistently outperforms vanilla RAG across all three backbones.

Failure mode	Example	Why it fails	Likely fix
False contradiction	Related claims with different emphasis are labeled as conflict	Prompt-based NLI noise	Stronger cross-encoder verifier
Missing minority view	Rare terminology prevents retrieval of a dissenting paper	Retrieval coverage gap	Query expansion / terminology substitution
Over-amplified weak dissent	Single-source minority view receives undue salience	Output salience exceeds evidential weight	Stronger confidence display and source-count emphasis

Table 10: Representative failure modes observed during manual inspection of low-CR queries.

Vanilla RAG	EVIRAG Full
<i>Query: Does homework improve academic achievement?</i>	
“Research consistently shows that homework improves academic achievement, with stronger effects for older students (grades 7–12). Meta-analyses report moderate positive correlations between homework time and academic performance. Homework builds self-regulation and reinforces classroom learning.” <i>CCS: 0.38. VC_{emb}: 0.33 (1/3 views covered).</i>	View 1 [Dominant, 6 sources, confidence: medium]: Conditional positive effect. Homework improves achievement for grades 7–12 ($d \approx 0.30$); effect near zero for primary grades. <i>CDA-7: methodological (grade-level scope differences).</i> View 2 [Alternative, 3 sources, confidence: medium]: Null meta-analytic effect when controlling for SES and teacher quality. <i>CDA-7: statistical + methodological.</i> View 3 [Minority, 2 sources, confidence: low]: High homework loads harm wellbeing with no academic benefit in high-achieving populations. <i>CDA-7: operational (achievement vs. wellbeing outcomes).</i> <i>Controversy class: Stable. CCS: 0.52. VC_{emb}: 1.00.</i>

Table 11: Homework query: Vanilla RAG collapses to the dominant conditional-benefit view; EVIRAG Full recovers all three ground-truth viewpoints with per-view source counts, CDA-7 causal attribution, and evidence-structured confidence.

Query	Vanilla RAG	EVIRAG Full (summary)
<i>Do statins prevent cardiovascular events in primary prevention?</i>	“Statins are recommended for high-risk primary prevention and reduce cardiovascular events.” (1 view; VC_{emb} : 0.33; CCS: 0.41)	View 1 [Dominant]: Benefit confirmed for high-risk groups (<i>CDA-7: population</i>). View 2: Marginal benefit with adverse-effect trade-off (<i>CDA-7: statistical</i>). View 3: Industry-bias concern undermining reported effect sizes (<i>CDA-7: methodological</i>). VC_{emb} : 1.00; Controversy class: Stable; CCS: 0.54.
<i>Does raising the minimum wage reduce employment?</i>	“Most studies find minimal employment effects from moderate minimum wage increases.” (1 view; VC_{emb} : 0.50; CCS: 0.44)	View 1 [Dominant]: Minimal aggregate employment effect (<i>CDA-7: methodological</i>). View 2: Significant job losses for low-wage workers in specific sectors (<i>CDA-7: population</i>). View 3: Positive employment via demand-side stimulus (<i>CDA-7: theoretical</i>). VC_{emb} : 1.00; Controversy class: Polarized; CCS: 0.50.

Table 12: Additional examples in biomedicine and economics. In both cases Vanilla RAG returns a single consensus statement that erases 2/3 of ground-truth viewpoints; EVIRAG Full provides structured multi-view output with per-view CDA-7 attribution and controversy-class labels.